

Page outlines: getting renderer data into a `chrome://` page


How `chrome://floating-window` shows the **level 1 and 2 headings** of every open tab, and why that one addition turned a synchronous page into an asynchronous one.

This is a companion to [floating-window-implementation.md](#), which covers the toolbar button, the bubble, and how the page is served. Read that first if you want the surrounding machinery; this document is only about the outlines, and goes deeper on the parts that are genuinely tricky. The options that were weighed and rejected are Part E of [floating-window-alternatives.md](#).

Source links. Every path below links into github.com/obeletski/chromium on the `floating-window` branch. Line anchors were checked against the tree these docs were written from; they will drift if the branch is rebased.

1. The asymmetry that shapes everything

The page shows two kinds of thing, and they could hardly be less alike.




The **tab list** — which windows exist, which tabs they hold, each tab's title, URL and index — is browser-process state. Reading it is a few pointer hops on the UI thread, and it cannot fail.


The **outline** of each page is DOM. It lives in whichever renderer process is hosting that tab, the browser holds no copy of it, and getting it means asking another process and waiting. It can be slow, it can be truncated, and it can never arrive at all.

Everything else in this document follows from that one difference.

2. Why the response became asynchronous

The page is served by a `WebUIDataSource` request filter, unchanged since the feature rendered one line of static text. What changed is *when* the filter can answer.






The crucial part is that this needed **no new mechanism**. `WebUIDataSource::GotDataCallback` was always allowed to be answered later — the API exists in callback form precisely because "a data source is allowed to answer asynchronously (reading from disk, waiting on a service)". The static page and the tab table simply happened to run it inline. Moving it into an object that runs it later is using the interface as designed, not working around it.

What this buys, and it is the whole reason the design holds together:

- no placeholder document and no second navigation — the renderer receives one complete page;
- no script in the served page, so the WebUI Content Security Policy never becomes a problem (see §7 of the implementation doc);
- no message handler and no `.mojom` of our own.

3. Assembling the page

This is the diagram to read if you read only one.



Mermaid aside, since this document is mostly diagrams. A `;` inside sequence-diagram `Note` or message text silently ends the statement — mermaid treats it as a line separator — and the rest of the note runs into the following lines until the parse fails somewhere unrelated. Quoted flowchart labels are unaffected. Worth knowing before debugging a "Parse error on line N" that points at an innocent line.


Why the sentinel exists

`pending_` starts at one *before* any snapshot is requested, and that extra count is dropped only after the request loop has finished.

Without it, consider a request that completes inline — a frame that is gone by the time the call is made, so the reply callback is destroyed immediately and the default fires synchronously. On the first tab of five, `pending_` would go `1 → 0`, `Finish()` would run, and the page would be published built from one tab while four requests were still being issued. The sentinel makes "still issuing" an explicit state that has to be left deliberately.


It also handles the degenerate case for free: a profile where no tab has a live renderer issues no requests at all, and the trailing `ResolveOne()` drives `pending_` to zero right there — the page is served immediately instead of waiting out a two-second deadline for replies that were never requested.

The collector's states



`Finish()` is reachable from two directions that know nothing about each other, so it must be idempotent. A late reply arriving after the deadline still calls `ResolveOne()`, still decrements, and may still reach `Finish()` — which returns immediately.

Who keeps the collector alive




Two independent things keep the object alive and either may outlive the other, which is exactly the case `base::RefCounted` is for: the outstanding reply callbacks, and the collector's own deadline timer. A `unique_ptr` owned by `HandleRequest()` would be destroyed when that function returns, which is before any reply arrives.

The timer holds a `base::Unretained(this)` and that is safe *because* it is a member: `OneShotTimer`'s destructor cancels it, so it cannot outlive the object it points at.

TabEntry holds strings, not `WebContents*`. A tab can be closed while its snapshot is in flight. A copied title and URL cannot dangle, so a tab that disappears mid-gather still renders as the row it was, rather than crashing or being silently dropped. This is a deliberate copy, not laziness.

4. Four failure modes, three mechanisms

This is where the design earns its keep, and where the obvious tool is not enough.



The red box is the point of this section.

`mojo::WrapCallbackWithDefaultInvokeIfNotRun ( callback\_helpers.h )` is the in-tree idiom for "this reply might not come", and it is what `ai\_data\_keyed\_service.cc` wraps around this very call. It is easy to reach for it and consider the problem solved. But it fires when the callback is **destroyed** — a renderer that is alive and spinning in a synchronous loop destroys nothing, holds the callback indefinitely, and is completely invisible to it.

Because the response is not composed until every tab has reported, a single wedged tab would leave the floating window blank for as long as it stayed wedged. The `kOverallDeadline` timer is the only thing standing between the feature and that outcome.


This was tested rather than argued: a page that busy-loops its main thread for sixty seconds. The window still opens, that tab reads *outline unavailable*, and every other tab renders its outline normally.

A note on what the deadline does not do. It does not cancel the outstanding request. The mojo callback stays alive until the renderer eventually replies or goes away; `Finish()` simply stops caring. The collector lives on until that last reference drops, and the late reply lands in a `Finished` state that ignores it.

5. Reading headings out of an accessibility tree

The outline is not obtained by running script in the page. It comes from


`WebContents::RequestAXTreeSnapshot() ( web\_contents.h )`, a one-shot capture that does not permanently change the accessibility mode.



An `AXTreeUpdate` (`ax_tree_update.h`) is a **flat** `std::vector<AXNodeData>`, not a tree — the structure lives in each node's `child_ids`. That is convenient here rather than annoying: the snapshotter serializes in document order, so reading `nodes` straight through yields headings in the order they appear on the page, which is exactly what an outline wants. No recursion, no tree reconstruction.

Trap 1 — the level is only serialized under `kExtendedProperties`


Blink writes `kHierarchicalLevel` inside a branch guarded on that AX mode flag (`ax_object.cc`, the `IsHeading(role) && HeadingLevel()` case). Ask for plain `kWebContents` and the snapshot still contains heading *nodes* — they simply all report level 0, so an h1/h2 filter drops every one of them. Nothing logs, nothing fails; the outline is just always empty.



The mode is spelled out as `kWebContents | kExtendedProperties` rather than reusing `ui::kAXModeWebContentsOnly` (`ax_mode.h`), which every in-tree snapshot caller passes — Compose, Read Anything, Mahi, `ai_data_keyed_service` — and which is those two flags **plus** `kInlineTextBoxes`. That third flag makes the renderer lay out and serialize per-word text boxes for the entire document. Nothing here reads them, and the cost would be paid once per open tab, every time the window is opened.

Trap 2 — "level 1 and 2" is not `querySelectorAll("h1, h2")`

Blink's `HeadingLevel()` (`ax_node_object.cc`) resolves the level like this:



Two consequences, both verified against real pages:

- `<div role="heading" aria-level="2">` **is** an h2 for this purpose, and an `<h1 aria-level="4">` is excluded. A CSS selector would get both backwards. This is the right answer — an outline is a question about document semantics, not about tag names.
- `GetComputedHeadingOffset()` is easy to over-read. It does **not** come from nesting inside `<section>` or `<article>`; it comes from the `headingoffset` and `headingreset` content attributes ([element.cc](#)), and it is gated on the `HeadingOffset` runtime feature, which is `status: "experimental"` and therefore off in a default build. On any normal page today, *level* is the tag number: an `<h1>` nested two `<section>` s deep still reports 1.

Trap 3 — an accessible name is not `textContent`

The name is *computed*, which cuts both ways:


- a heading whose markup wraps across several source lines carries those newlines and their indentation, so names go through `base::CollapseWhitespaceASCII()` ([string_util.h](#));
- a heading containing a `<script>` element contributes none of that element's text, which is a pleasant accident — `<h2>Costs & <script>...</script> risks</h2>` comes out as `Costs & risks`.

Names are then capped with `base::TruncateUTF8ToByteSize()`, which cuts on a character boundary where a naive `substr()` would split a multi-byte sequence and emit invalid UTF-8.

Snapshots are requested with `kSameOriginDirectDescendants`: a cross-origin iframe's headings are not part of this page's outline, and reaching into them would widen what a `chrome://` page reads out of arbitrary sites for no benefit.

6. What ends up in the document

The outline rows are rows of the same table as the tabs, spanning the title and URL columns.



(The arrows from `tr.tab` are "followed by", not nesting — HTML tables have no row nesting. The visual hierarchy is entirely the `lvl1` / `lvl2` indent.)

Abridged from a real run:

```

<tr class="tab" aria-current="true"><td class="idx">1</td><td>Quarterly report – draft</td><td>
<tr class="hd lvl1"><td class="idx"></td><td colspan="2">Quarterly Report</td></tr>
<tr class="hd lvl2"><td class="idx"></td><td colspan="2">Revenue</td></tr>
<tr class="hd lvl2"><td class="idx"></td><td colspan="2">Costs & risks</td></tr>
<tr class="hd lvl1"><td class="idx"></td><td colspan="2">Appendix</td></tr>

```

Three decisions worth naming:

- **Rows, not a nested `<ul>` in the title cell.** A list would have to opt out of the `nowrap` / `text-overflow` rules the cells rely on, and it would break the alignment of the index column.
- **The indent is on the cell, not the row.** Each outline row still carries an empty `<td class="idx">` so the tab numbering column stays aligned; padding goes on `td:last-child`.
- **Two different empty states.** "no level 1 or 2 headings" means the snapshot came back and the page genuinely has none. "outline unavailable" means it never came back. Collapsing them into one message would hide the difference between "this page has no structure" and "we could not read it".

7. Costs, caps, and what is deliberately missing

Knob	Value	Why
<code>kOverallDeadline</code>	2 s	the browser-side backstop; the only cover for a live-but-silent renderer
<code>kSnapshotTimeout</code>	1.2 s	passed to the renderer, which truncates the tree it serializes
<code>kMaxAxNodesPerTab</code>	20 000	a cap on a hostile or enormous document
<code>kMaxHeadingsPerTab</code>	12	beyond this the row reads + N more
<code>kMaxHeadingBytes</code>	300	per heading, cut on a UTF-8 boundary

Known limits, none of them accidental:

- **One IPC per tab, per opening.** Nothing is cached between openings and nothing is batched. With many tabs the window appears measurably later than it did when the page was pure browser-process state. This is the real price of the feature.
- **Levels 1 and 2 only.** Deeper headings are dropped rather than folded in.
- **Cross-origin iframes contribute nothing**, by choice of snapshot policy. A page whose real content sits in a cross-origin iframe will look emptier than it is.
- **Still a snapshot.** Like the tab list, the outlines are read once while the response is composed. The bubble builds a fresh `WebContents` on every press, so reopening is the refresh gesture.

8. Verification

Driven under `xvfb`, with real X input synthesised through `libXtst` via `ctypes` to click the toolbar button, plus CDP `Page.navigate` to load `chrome://floating-window` into an ordinary tab for markup inspection. (`/json/new` refuses `chrome://` URLs and `--headless` accepts only one target, so neither is a route to the page.)

Case	Expected	Result
h1 , h2 , h3 , h2 , h1	the h3 excluded, four rows in document order	☐
<h1> nested one and two <section> s deep	all level 1	☐
<div role="heading" aria-level="2">	included as level 2	☐
<div role="heading" aria-level="4">	excluded	☐
<h2> split across three source lines	collapsed to one line	☐
heading containing <script>alert(1) </script>	script text absent, rest escaped	☐
page with no headings	<i>no level 1 or 2 headings</i>	☐
renderer busy-looping for 60 s	window still opens; that tab <i>outline unavailable</i> ; others normal	☐

No CSP violations, FATAL or DCHECK output in the browser log across any run. gn check passes on //chrome/browser/ui/webui/floating_window .

The last row is the one worth having actually run: it is the failure that the first-choice idiom does not cover, and without the deadline the window would have stayed blank indefinitely.